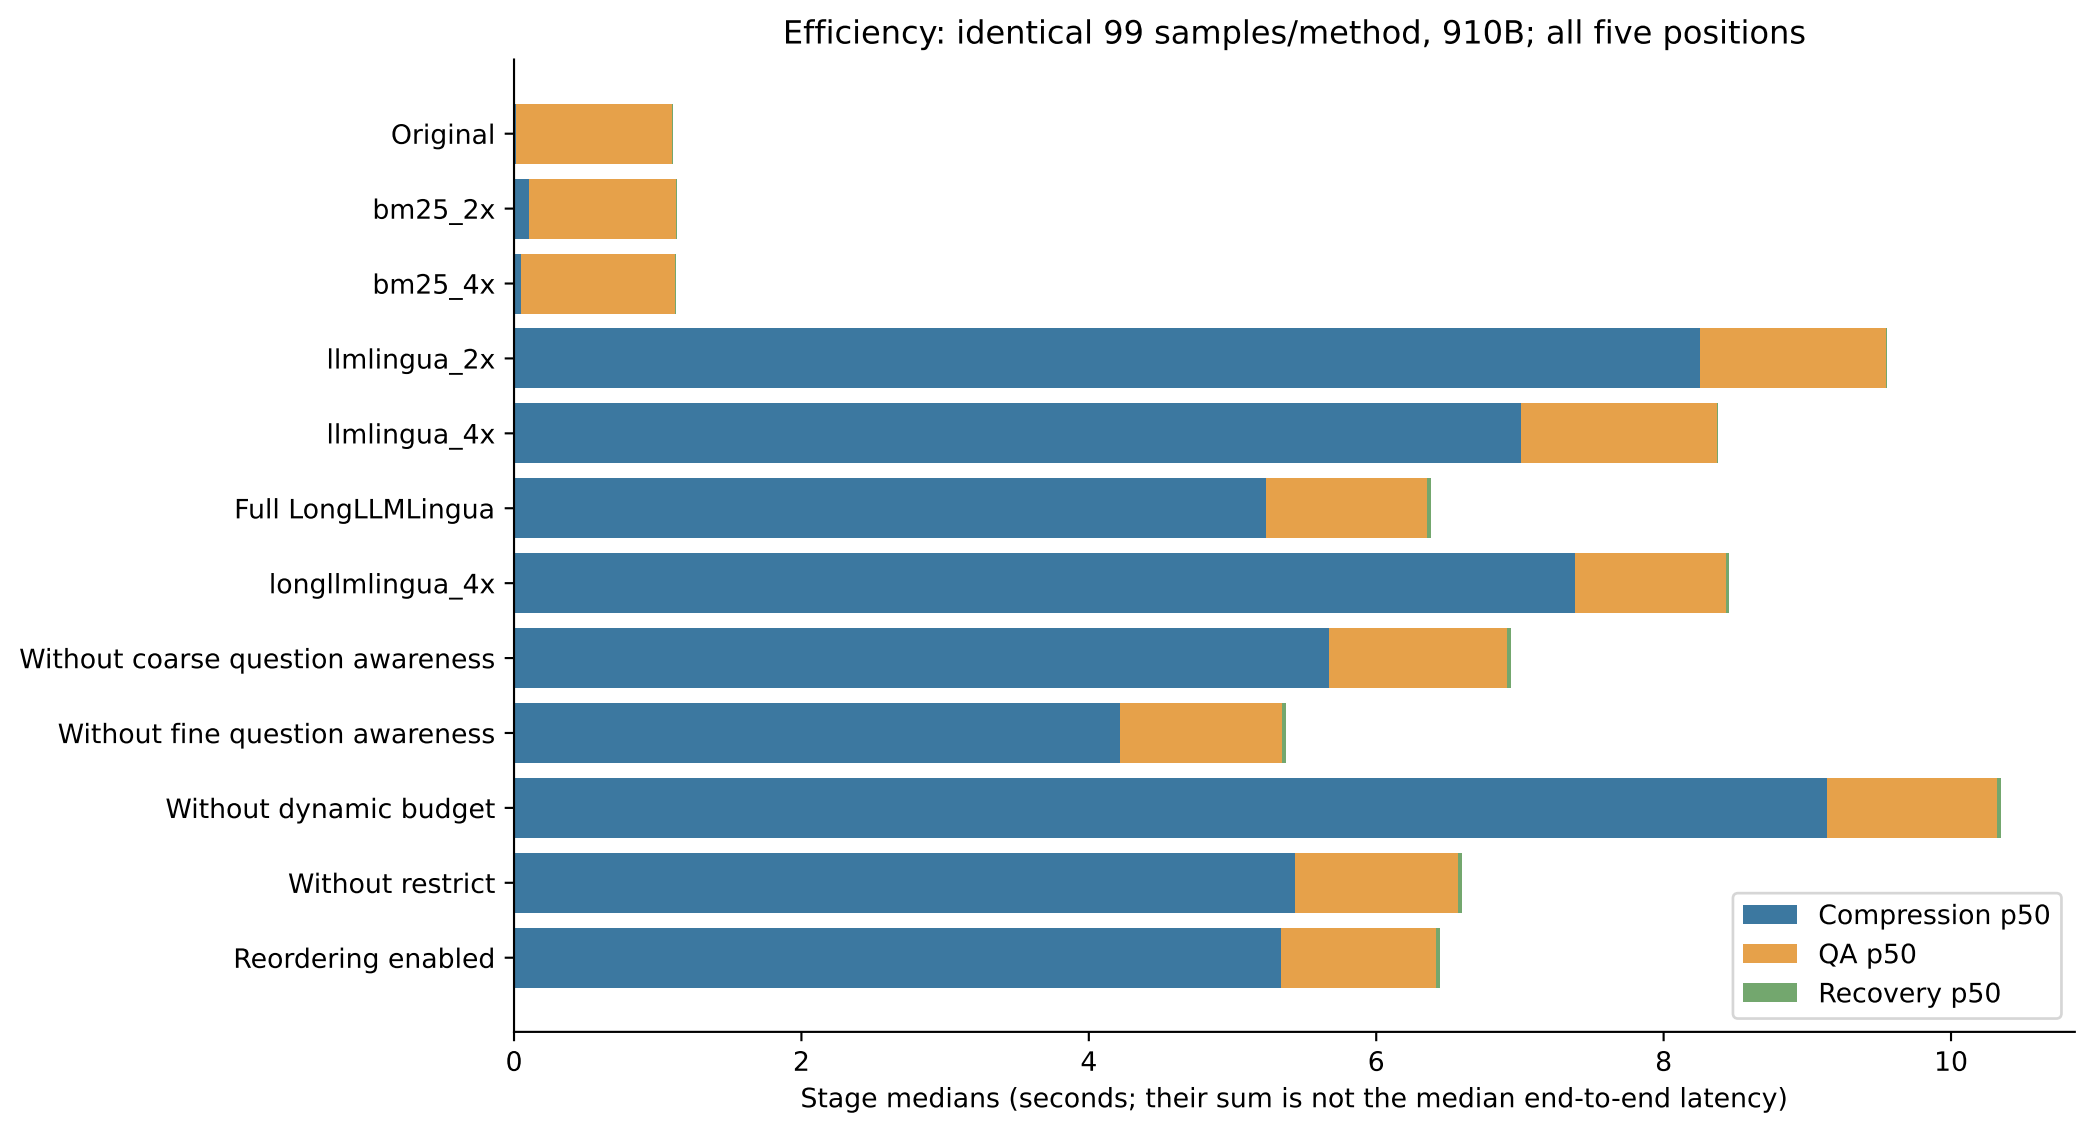
Efficiency: identical 99 samples/method, 910B; all five positions
Original
bm25_2x
bm25_4x
llmlingua_2x
llmlingua_4x
Full LongLLMLingua
longllmlingua_4x
Without coarse question awareness
Without fine question awareness
Without dynamic budget
Without restrict
Reordering enabled
Compression p50
QA p50
Recovery p50
0
2
4
6
8
10
Stage medians (seconds; their sum is not the median end-to-end latency)